

HW 5 (I have included my comments in # and my code)

Question 11.1

Using the crime data set `uscrime.txt` from Questions 8.2, 9.1, and 10.1, build a regression model using:

1. Stepwise regression
2. Lasso
3. Elastic net

```
uscrime <- read.table("11.uscrimeSummer2018.txt", stringsAsFactors = FALSE, header = TRUE)
```

```
#Stepwise Regression
```

```
#Start with 15 predictors and work its way down. It starts with 15 and reduces to 8 and uses AIC to do that.
```

```
model_1 <- lm(Crime ~ ., data = uscrime)
```

```
stepAIC(model_1, direction = "backward")
```

```
Coefficients:
(Intercept)      M      Ed      Po1      M.F      U1      U2      Ineq      Prob
-6426.10      93.32     180.12    102.65     22.34    -6086.63     187.35     61.33    -3796.03
```

```
#Use the 8 factors to run a regression.
```

```
model_2 <- lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob, data = uscrime)
```

```
summary(model_2)
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -6426.10    1194.61   -5.379  4.04e-06 ***
M              93.32      33.50    2.786  0.00828 **
Ed            180.12      52.75    3.414  0.00153 **
Po1           102.65      15.52    6.613  8.26e-08 ***
M.F            22.34      13.60    1.642  0.10874
U1           -6086.63   3339.27   -1.823  0.07622 .
U2            187.35      72.48    2.585  0.01371 *
Ineq           61.33      13.96    4.394  8.63e-05 ***
Prob          -3796.03   1490.65   -2.547  0.01505 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 195.5 on 38 degrees of freedom
Multiple R-squared:  0.7888,    Adjusted R-squared:  0.7444
F-statistic: 17.74 on 8 and 38 DF,  p-value: 1.159e-10
```

```
#Remove M.F & U1 since it is insignificant
```

```
model_2 <- lm(formula = Crime ~ M + Ed + Po1 + U2 + Ineq + Prob, data = uscrime)
```

```
summary(model_2)
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -5040.50     899.84   -5.602  1.72e-06 ***
M              105.02      33.30    3.154  0.00305 **
Ed            196.47      44.75    4.390  8.07e-05 ***
Po1           115.02      13.75    8.363  2.56e-10 ***
U2             89.37      40.91    2.185  0.03483 *
Ineq           67.65      13.94    4.855  1.88e-05 ***
Prob          -3801.84   1528.10   -2.488  0.01711 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 200.7 on 40 degrees of freedom
Multiple R-squared:  0.7659,    Adjusted R-squared:  0.7307
F-statistic: 21.81 on 6 and 40 DF,  p-value: 3.418e-11
```

```
install.packages("glmnet") # for LASSO and Elastic net
library(glmnet)
set.seed(42)
datascale<-scale(uscrime)
#Lasso
model_lasso <- cv.glmnet(x=as.matrix(uscrime[,-16]), y=as.matrix(uscrime[,16]), alpha = 1, nfolds = 5,
type.measure = "mse", family = "gaussian")
model_lasso

$lambda.min
[1] 17.71724

$lambda.1se
[1] 49.29927

attr(,"class")
[1] "cv.glmnet"
```

model_lasso\$lambda #lambda is the t values or budgets.We want to pick t value that gives lowest error.

model_lasso\$cvm

```
> model_lasso$lambda #lambda is the t values or budgets.We want to pick t value that gives lowest error.
[1] 263.0953966 239.7227267 218.4264204 199.0220193 181.3414516 165.2315769 150.5528590 137.1781579 124.9916285 113.8877167
[11] 103.7702459 94.5515832 86.1518812 78.4983855 71.5248053 65.1707387 59.3811499 54.1058922 49.2992739 44.9196623
[21] 40.9291233 37.2930928 33.9800772 30.9613808 28.2108571 25.7046823 23.4211492 21.3404788 19.4446495 17.7172404
[31] 16.1432896 14.7091643 13.4024427 12.2118066 11.1269433 10.1384564 9.2377838 8.4171246 7.6693704 6.9880447
[41] 6.3672461 5.8015975 5.2861996 4.8165882 4.3886957 3.9988161 3.6435723 3.3198874 3.0249577 2.7562288
[51] 2.5113731 2.2882696 2.0849860 1.8997616 1.7309920 1.5772155 1.4371000 1.3094320 1.1931057 1.0871134
[61] 0.9905373 0.9025407 0.8223615 0.7493051 0.6827389 0.6220863 0.5668219 0.5164670 0.4705855 0.4287799
> model_lasso$cvm
[1] 146403.92 140712.28 135089.27 128984.80 121931.30 116091.21 111257.10 107256.90 103947.86 101211.55 99056.47 97443.86
[13] 96159.04 94984.45 93391.12 91357.34 88477.78 84222.76 79891.03 76096.80 73026.20 70760.11 69371.69 68695.59
[25] 68337.95 68270.79 68340.70 68161.82 67682.72 67481.51 67707.48 68263.88 68747.49 69116.26 69787.49 70733.14
[37] 71710.68 72718.18 73796.45 74940.75 76101.91 77328.42 78578.87 79768.80 80896.15 81926.74 82855.69 83845.20
[49] 84866.33 85884.60 86902.47 87842.83 88717.25 89537.06 90297.60 91102.65 92089.00 93167.84 94199.94 95169.05
[61] 96102.33 97024.02 97908.63 98754.59 99550.90 100291.01 100975.68 101610.07 102143.30 102609.97
```

coef(model_lasso, s= model_lasso\$lambda.min) #this will have the lowest mean square errors

```
(Intercept) -3828.8353017
M            56.1008808
So           30.7597658
Ed           70.8167194
Po1          103.2100909
Po2          .
LF           .
M.F          16.7898439
Pop          .
NW           0.3226147
U1           .
U2           24.9099830
Wealth       .
Ineq         37.7315902
Prob         -3179.3760049
Time         .
```

```
model_3 <-lm(formula = Crime~ M+So+Ed+Po1+M.F+Pop+NW+U1+U2+Wealth+Ineq+Prob, data =
uscrime)
summary(model_3)
```

```

(Intercept) -6.393e+03  1.413e+03 -4.524 7.05e-05 ***
M            8.968e+01  3.927e+01  2.284 0.02876 *
So           2.289e+01  1.253e+02  0.183 0.85621
Ed           1.749e+02  5.627e+01  3.109 0.00378 **
Po1          9.865e+01  2.187e+01  4.511 7.32e-05 ***
M.F          1.660e+01  1.633e+01  1.017 0.31656
Pop          -8.734e-01  1.199e+00 -0.729 0.47113
NW           1.863e+00  5.613e+00  0.332 0.74195
U1           -4.979e+03  3.643e+03 -1.367 0.18069
U2           1.667e+02  7.906e+01  2.108 0.04245 *
wealth       8.633e-02  9.900e-02  0.872 0.38932
Ineq         7.163e+01  2.135e+01  3.355 0.00196 **
Prob        -4.079e+03  1.809e+03 -2.255 0.03065 *
---
signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 202.6 on 34 degrees of freedom
Multiple R-squared:  0.7971,    Adjusted R-squared:  0.7255
F-statistic: 11.13 on 12 and 34 DF,  p-value: 1.52e-08

```

#Elastic

```

model_Elastic <- cv.glmnet(x=as.matrix(uscrime[, -16]), y=as.matrix(uscrime[, 16]), alpha = .5, nfolds = 5,
type.measure = "mse", family = "gaussian")
model_Elastic

```

```

$lambda.min
[1] 11.60319

$lambda.1se
[1] 42.68096

attr(,"class")
[1] "cv.glmnet"

```

model_Elastic\$lambda #lambda is the t values or budgets. We want to pick t value that gives lowest error.

model_Elastic\$cvm

```

> model_Elastic$lambda #lambda is the t values or budgets. We want to pick t value that gives lowest error.
[1] 526.1907933 479.4454535 436.8528408 398.0440385 362.6829032 330.4631537 301.1057180 274.3563159 249.9832570 227.7754334
[11] 207.5404917 189.1031664 172.3037623 156.9967710 143.0496106 130.3414774 118.7622998 108.2117845 98.5985478 89.8393245
[21] 81.8582466 74.5861856 67.9601544 61.9227616 56.4217141 51.4093646 46.8422983 42.6809576 38.8892990 35.4344809
[31] 32.2865792 29.4183285 26.8048853 24.4236132 22.2538867 20.2769127 18.4755677 16.8342492 15.3387409 13.9760894
[41] 12.7344922 11.6031950 10.5723991 9.6331763 8.7773915 7.9976322 7.2871446 6.6397748 6.0499155 5.5124577
[51] 5.0227461 4.5765392 4.1699721 3.7995232 3.4619841 3.1544309 2.8742000 2.6188640 2.3862113 2.1742269
[61] 1.9810746 1.8050814 1.6447229 1.4986103 1.3654779 1.2441726 1.1336437 1.0329339 0.9411709 0.8575599
[71] 0.7813766 0.7119613 0.6487126 0.5910828 0.5385726 0.4907273 0.4471324 0.4074104

> model_Elastic$cvm
[1] 149625.72 146235.01 141410.17 134461.08 128124.68 122588.08 117824.93 113745.75 110354.42 107793.67 105983.63 104743.87
[13] 103797.11 103126.45 102736.90 102287.92 100683.61 98511.04 96376.61 93940.52 91400.46 88893.98 86330.24 83963.75
[25] 81468.88 79086.75 76742.70 74249.60 72124.56 70391.48 68848.49 67510.54 66409.80 65511.68 64741.12 64029.00
[37] 63443.45 62950.03 62577.11 62333.41 62185.51 62127.22 62171.64 62238.60 62409.74 62671.18 62948.48 63235.78
[49] 63552.84 63910.79 64322.79 64743.39 65168.30 65578.49 65962.82 66329.33 66672.59 67011.38 67372.25 67689.96
[61] 67990.89 68250.50 68409.98 68546.12 68684.21 68824.65 68966.41 69109.63 69256.49 69505.09 69693.50 69835.96
[73] 69975.43 70110.07 70238.25 70353.65 70466.05 70573.11

```

coef(model_Elastic, s= model_Elastic\$lambda.min) #this will have the lowest mean square errors

```

(Intercept) -5.448218e+03
M            7.317505e+01
So           4.726900e+01
Ed           1.285266e+02
Po1          8.286051e+01
Po2          1.379667e+01
LF           2.007735e+01
M.F          2.139471e+01
Pop          -2.114354e-03
NW           1.599082e+00
U1           -3.038527e+03
U2           1.078048e+02
wealth       3.039998e-02
Ineq         4.969286e+01
Prob        -3.823196e+03
Time         .

```

```

model_4 <- lm(formula = Crime ~ M + So + Ed + Po1 + Po2 + M.F + Pop + NW + U1 + U2 + Wealth + Ineq + Prob, data =
uscrime)
summary(model_4)

```

```

M      8.743e+01  3.964e+01  2.205  0.034514 *
So      3.440e+01  1.271e+02  0.271  0.788398
Ed      1.809e+02  5.721e+01  3.163  0.003346 **
Po1     1.688e+02  9.667e+01  1.746  0.090115 .
Po2     -7.692e+01  1.032e+02 -0.745  0.461484
M.F     1.474e+01  1.663e+01  0.887  0.381622
Pop     -9.510e-01  1.211e+00 -0.785  0.437837
NW      2.422e+00  5.699e+00  0.425  0.673604
U1      -4.805e+03  3.674e+03 -1.308  0.200017
U2      1.622e+02  7.982e+01  2.032  0.050269 .
wealth  8.501e-02  9.967e-02  0.853  0.399833
Ineq    6.912e+01  2.175e+01  3.177  0.003219 **
Prob    -4.185e+03  1.826e+03 -2.292  0.028430 *
---
signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 204 on 33 degrees of freedom
Multiple R-squared:  0.8005,    Adjusted R-squared:  0.7219
F-statistic: 10.19 on 13 and 33 DF,  p-value: 4.088e-08

```

Question 12.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a design of experiments approach would be appropriate.

For example, let's say that I want to introduce a new organic vegan cereal product to the market. I try to create different kinds of packaging for my cereal to get an idea of which packaging would be appealing to the consumers. I can use DOE to understand the effect of packaging on a set of consumers to understand their buying behavior.

Question 12.2

To determine the value of 10 different yes/no features to the market value of a house (large yard, solar roof, etc.), a real estate agent plans to survey 50 potential buyers, showing a fictitious house with different combinations of features. To reduce the survey size, the agent wants to show just 16 fictitious houses. Use R's `FrF2` function (in the `FrF2` package) to find a fractional factorial design for this experiment: what set of features should each of the 16 fictitious houses have? Note: the output of `FrF2` is "1" (include) or "-1" (don't include) for each feature.

#Need two inputs (nruns = 16 fictitious houses and nfactors = 10 features). Below is the result of the fractional factorial design.

`FrF2(16, 10)`

```

> FrF2(16, 10)
      A  B  C  D  E  F  G  H  J  K
1      1  1 -1 -1  1 -1 -1 -1  1  1
2     -1 -1 -1 -1  1  1  1  1 -1  1
3     -1  1  1  1 -1 -1  1 -1  1 -1
4      1 -1 -1 -1 -1 -1  1 -1 -1 -1
5     -1 -1  1  1  1 -1 -1 -1 -1  1
6      1 -1 -1  1 -1 -1  1  1  1  1
7      1 -1  1 -1 -1  1 -1 -1  1  1
8     -1 -1 -1  1  1  1  1 -1  1 -1
9      1 -1  1  1 -1  1 -1  1 -1 -1
10     -1  1  1 -1 -1 -1  1  1 -1  1
11     -1  1 -1  1 -1  1 -1 -1 -1  1
12     -1  1 -1 -1 -1  1 -1  1  1 -1
13      1  1  1 -1  1  1  1 -1 -1 -1
14      1  1  1  1  1  1  1  1  1  1
15     -1 -1  1 -1  1 -1 -1  1  1 -1
16      1  1 -1  1  1 -1 -1  1 -1 -1
class=design, type= FrF2

```

Question 13.1

For each of the following distributions, give an example of data that you would expect to follow this distribution (besides the examples already discussed in class).

- a. Binomial : The binomial is a type of distribution that has two possible outcomes. Data on students overall grade in a course explained as either a pass or a fail.
- b. Geometric: you ask people outside a polling station who they voted for until you find someone that voted for the independent candidate in a local election. The geometric distribution would represent the number of people who you had to poll before you found someone who voted independent.
- c. Poisson: Given the number of diners in a certain restaurant every day, if the average number of diners for seven days is 500, you can predict the probability of a certain day having more customers.
- d. Exponential: Let's say a Poisson distribution models the number of births in a given time period. The time in between each birth can be modeled with an exponential distribution
- e. Weibull: How long will it take for a TV to become defective since the time it has been switched on.